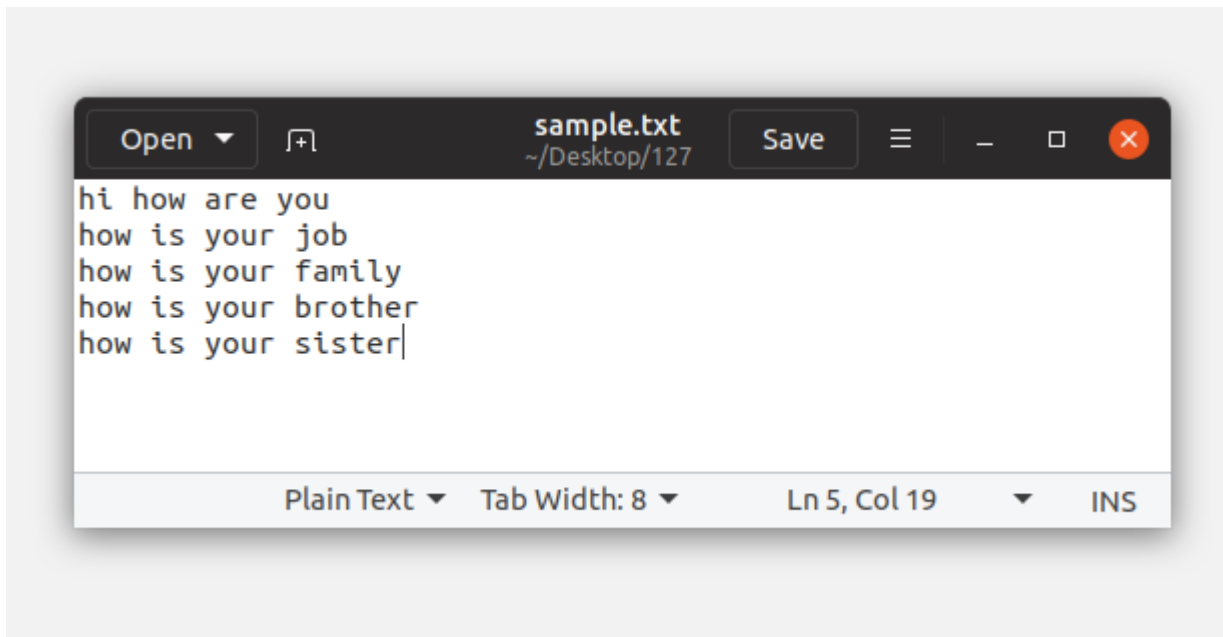


Screen Shots of Execution

```
hduser@bmsce-Precision-T1700: ~$ cd / Desktop / 127
hduser@bmsce-Precision-T1700:~$ start-all.sh
This script is deprecated. Instead use start-dfs.sh and start-yarn.sh
Starting namenodes on [localhost]
hduser@localhost's password:
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-hduser-namenode-bmsce-Precision-T1700.out
hduser@localhost's password:
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hduser-datanode-bmsce-Precision-T1700.out
Starting secondary namenodes [0.0.0.0]
hduser@0.0.0.0's password:
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-hduser-secondarynamenode-bmsce-Precision-T1700.out
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-hduser-resourcemanager-bmsce-Precision-T1700.out
hduser@localhost's password:
localhost: starting nodenanager, logging to /usr/local/hadoop/logs/yarn-hduser-nodenanager-bmsce-Precision-T1700.out
hduser@bmsce-Precision-T1700:~$ jps
6512 DataNode
7217 NodeManager
6884 ResourceManager
6725 SecondaryNameNode
6344 NameNode
7339 Jps
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /
Found 3 items
drwxr-xr-x - hduser supergroup          0 2019-10-23 15:08 /dg
drwxrwxr-x - hduser supergroup          0 2019-10-23 15:53 /tmp
drwxr-xr-x - hduser supergroup          0 2019-08-01 16:03 /user
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /
Found 3 items
drwxr-xr-x - hduser supergroup          0 2019-10-23 15:08 /dg
drwxrwxr-x - hduser supergroup          0 2019-10-23 15:53 /tmp
drwxr-xr-x - hduser supergroup          0 2019-08-01 16:03 /user
hduser@bmsce-Precision-T1700:~$ hadoop fs -mkdir /rgs
hduser@bmsce-Precision-T1700:~$ hadoop fs -copyFromLocal /home/bmsce/Desktop/127/sample.txt /rgs/test.txt
hduser@bmsce-Precision-T1700:~$ hadoop jar /home/bmsce/Desktop/127/WordCount.jar WCDriver input output
22/06/14 10:00:28 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
22/06/14 10:00:28 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
22/06/14 10:00:28 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with processName=JobTracker, sessionId = already initialized
22/06/14 10:00:28 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
Exception in thread "main" org.apache.hadoop.mapred.InvalidInputException: Input path does not exist: hdfs://localhost:54310/user/hduser/input
at org.apache.hadoop.mapred.FileInputFormat.singleThreadedListStatus(FileInputFormat.java:285)
at org.apache.hadoop.mapred.FileInputFormat.listStatus(FileInputFormat.java:226)
at org.apache.hadoop.mapred.FileInputFormat.getSplits(FileInputFormat.java:313)
at org.apache.hadoop.mapreduce.JobSubmitter.writeOldSplits(JobSubmitter.java:624)
at org.apache.hadoop.mapreduce.JobSubmitter.writeSplits(JobSubmitter.java:616)
at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:492)
at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1296)
at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1293)
at java.security.AccessController.doPrivileged(Native Method)
at javax.security.auth.Subject.doAs(Subject.java:422)
at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628)
at org.apache.hadoop.mapreduce.Job.submit(Job.java:1293)
at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:562)
at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:557)

at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:557)
at java.security.AccessController.doPrivileged(Native Method)
at javax.security.auth.Subject.doAs(Subject.java:422)
at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628)
at org.apache.hadoop.mapred.JobClient.submitJobInternal(JobClient.java:557)
at org.apache.hadoop.mapred.JobClient.submitJob(JobClient.java:548)
at org.apache.hadoop.mapred.JobClient.runJob(JobClient.java:833)
at WCDriver.run(WCDriver.java:33)
at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:79)
at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:84)
at WCDriver.main(WCDriver.java:40)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /output
ls: '/output': No such file or directory
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /
Found 4 items
drwxr-xr-x - hduser supergroup          0 2019-10-23 15:08 /dg
drwxr-xr-x - hduser supergroup          0 2022-06-14 09:58 /rgs
drwxrwxr-x - hduser supergroup          0 2019-10-23 15:53 /tmp
drwxr-xr-x - hduser supergroup          0 2019-08-01 16:03 /user
hduser@bmsce-Precision-T1700:~$ hadoop fs -cat /output/part-00000
cat: '/output/part-00000': No such file or directory
hduser@bmsce-Precision-T1700:~$ cd rgs
bash: cd: rgs: No such file or directory
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /
Found 4 items
drwxr-xr-x - hduser supergroup          0 2019-10-23 15:08 /dg
drwxr-xr-x - hduser supergroup          0 2022-06-14 09:58 /rgs
drwxrwxr-x - hduser supergroup          0 2019-10-23 15:53 /tmp
drwxr-xr-x - hduser supergroup          0 2019-08-01 16:03 /user
hduser@bmsce-Precision-T1700:~$ cd /rgs
bash: cd: /rgs: No such file or directory
hduser@bmsce-Precision-T1700:~$ hdfs dfs /cd rgs
/cd: Unknown command
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cd /rgs
-cd: Unknown command
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cd /rgs/
-cd: Unknown command
hduser@bmsce-Precision-T1700:~$ hadoop fs -cd rgs
-cd: Unknown command
hduser@bmsce-Precision-T1700:~$ hdfs dfs -ls /rgs
Found 1 items
-rw-r--r-- 1 hduser supergroup          89 2022-06-14 09:58 /rgs/test.txt
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /rgs/test.txt
hi how are you
how is your job
how is your family
how is your brother
hduser@bmsce-Precision-T1700:~$ hadoop fs -copyFromLocal /home/bmsce/Desktop/hduser@bmsce-Precision-T1700hduser@bmsce-Precision-T1700hduser@bmsce-Precision-T1700hduser@bmsce-Precision-T1700hduser@bmsce-Precision-T1700-S hadoo
```


Input Sample file :



Final Output :

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /output/*  
are      1  
brother  1  
family   1  
hi        1  
how       5  
is        4  
job       1  
sister    1  
you       1  
your      4  
hduser@bmsce-Precision-T1700:~$
```

Program :

Create Three Java Classes into the project. Name them WCDriver(having the main function), WCMapper, WCReducer.

```
// Importing libraries
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

public class WCMapper extends MapReduceBase implements Mapper<LongWritable,Text,
Text, IntWritable> {

    // Map function
    public void map(LongWritable key, Text value, OutputCollector<Text,
        IntWritable> output, Reporter rep) throws IOException
    {

        String line = value.toString();

        // Splitting the line on spaces
        for (String word : line.split(" "))
        {
            if (word.length() > 0)
            {
                output.collect(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

Reducer Code: You have to copy paste this program into the WCReducer Java Class file.

// Importing libraries

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
```

```
public class WCReducer extends MapReduceBase implements Reducer<Text,IntWritable,
Text, IntWritable> {
```

```
    // Reduce function
```

```
    public void reduce(Text key, Iterator<IntWritable> value,
                        OutputCollector<Text, IntWritable> output,
                        Reporter rep) throws IOException
```

```
    {
```

```
        int count = 0;
```

```
        // Counting the frequency of each words
```

```
        while (value.hasNext())
```

```
        {
```

```
            IntWritable i = value.next();
```

```
            count += i.get();
```

```
        }
```

```
        output.collect(key, new IntWritable(count));
```

```
    }
```

```
}
```

Driver Code: You have to copy paste this program into the WCDriver Java Class file.

// Importing libraries

```
import java.io.IOException;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
```

```
public class WCDriver extends Configured implements Tool {
```

```
    public int run(String args[]) throws IOException
    {
```

```
        if (args.length < 2)
        {
            System.out.println("Please give valid inputs");
            return -1;
        }
```

```
        JobConf conf = new JobConf(WCDriver.class);
        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));
        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class);
        conf.setMapOutputKeyClass(Text.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        JobClient.runJob(conf);
        return 0;
    }
```

```
// Main Method
```

```
public static void main(String args[]) throws Exception
{
```

```
    int exitCode = ToolRunner.run(new WCDriver(), args);
    System.out.println(exitCode);
}
```

```
}
```

Export as jar file

Mapreduce Program

1. hduser@bmsce-Precision-T1700:/\$ su hduser

Password:

hduser@bmsce-Precision-T1700:/\$ cd /

hduser@bmsce-Precision-T1700:/\$ cd /usr/local/hadoop/sbin

2. hduser@bmsce-Precision-T1700:/usr/local/hadoop/sbin\$ start-all.sh

3. hduser@bmsce-Precision-T1700:/usr/local/hadoop/sbin\$ jps

4. create a file on Desktop(sample.txt) and type the below lines:

hi how are you

how is your job

how is your family

how is your brother

how is your sister

save your file

5. View the directory content

hadoop fs -ls /

6. Create a directory using the following command. If any directory existing, use the same directory for the command

hadoop fs -mkdir /rgs

7. Copy the file into HDFS

hadoop fs -copyFromLocal D:/sample.txt /rgs/test.txt

8.Run the Map Reduce Program

hadoop fs -ls /output/

hadoop fs -copyFromLocal /home/hduser/Desktop/file1.txt /rgs/test.txt

hadoop jar /home/hduser/Desktop/MapReduceClient.jar WCDriver /rgs/test.txt /output/

hdfs dfs -cat /output/*

My File

